

Where Liquid Networks Help, Hurt, and Fail to Train: A Pre-Registered Boundary Map

Olga Żydziaś
Independent Research

2026-07-24

Abstract

Liquid neural networks — continuous-time cores such as LTCs and their closed-form successor (CfC), wired with NCP/AutoNCP — are claimed to grant perception-and-control policies greater robustness under distribution shift. We subject that claim to a pre-registered measurement program spanning three task regimes at core-parameter parity and equal training budget: open-loop classification (companion P0, PASTIS), open-loop onset detection (companion E6), and closed-loop control (state-loop, LiquidFlight v1.0; vision-loop, this work). In the richest regime — a vision-in-the-loop fly-to-target twin with a perceptual-shift ladder — no CfC arm reaches our pre-registered nominal-capability precondition (mean $\geq 90\%$ success at core parity of $27,648 \pm 2\%$) on either allowed learning rate, so the robustness claim is *untestable under our precondition* and we ran *zero* out-of-distribution (OOD) evaluations. Beneath that boundary we measure a *trainability tax*: at identical budget, a GRU control reaches 100% while the wired CfC peaks at 92% (per-leg means up to 75.3%, a 43 pp seed spread) and the dense CfC caps at 65%. We attribute trainability to wiring, localize a continuous-time Δt implementation dependency, and contribute a reusable gate/precondition methodology with an arithmetic early-stopping rule. A negative result, pre-registered, maps where the mechanism can and cannot be measured.

1 Introduction

Continuous-time recurrent networks have become an attractive substrate for embodied perception and control. Liquid Time-Constant networks (LTCs) and their closed-form counterpart (CfC), especially when wired with Neural Circuit Policies (NCP/AutoNCP), have been reported to yield policies that generalize better under distribution shift than conventional gated recurrences such as GRUs. The most cited demonstration in the flight domain reports that a compact liquid policy attends to the task-relevant part of the visual scene and holds up under perceptual perturbation [Chahine et al., 2023].

Two properties of that literature motivate this work. First, demonstrations are typically reported **without core-parameter parity** against a matched baseline and **without pre-registration** of the success criterion; a favorable curve is shown, but the counterfactual — a same-budget, same-data baseline under a criterion fixed in advance — is often absent. Second, the field does not report a **trainability layer**: whether the liquid core can be brought to competence at all, at parity, under a fixed budget, before any robustness question is even asked.

We take the claim seriously enough to try to falsify it under conditions it should survive. We build an honest twin — identical encoder, identical head, identical data and budget, cores matched to within $\pm 2\%$ — and we freeze the success criterion before measuring. This paper reports the vision-loop regime (phase 3) as the primary result and folds it into a portfolio-level **boundary map** with two companion regimes.

Contributions (mapped to our claim register C1–C8):

- A pre-registered vision-loop twin in which the robustness claim is *untestable under our pre-condition* — the deciding arm never qualifies, and we hold to zero OOD evaluations (C1).
- A measured **trainability tax** at parity: GRU 100% vs. wired-CfC (peak 92%, per-leg means up to 75.3%) vs. dense-CfC (cap 65%), at identical budget (C2).
- The finding that the boundary is **variance, not mean** — seed instability is the object of measurement, echoing the companion P0 population spread (C3).
- A **wiring attribution**: AutoNCP makes the CfC trainable where the dense CfC caps (C4).
- Engineering findings on the continuous-time Δt path and DAgger procedure (C5, C6).
- A transferable **methodology**: gates, preconditions, annexes with a stop rule, and arithmetic early resolution (C7); plus an instrument characterization (C8).

2 Related Work

Liquid networks. Liquid Time-Constant networks model the hidden state as a continuous-time dynamical system whose time constants are themselves input-dependent, giving the network a natively continuous notion of elapsed time [Hasani et al., 2021]. Their closed-form successor, the CfC, replaces the numerical ODE solve with a closed-form update that preserves the continuous-time gate while removing solver cost [Hasani et al., 2022]. Neural Circuit Policies and their automatic wiring (AutoNCP) impose a sparse, structured connectivity on the recurrent core, inspired by biological circuits and reported to improve auditability and compactness [Lechner et al., 2020]. The robustness claim we test derives largely from a flight-navigation demonstration in which a small wired liquid policy generalizes out of distribution and attends to task-relevant structure [Chahine et al., 2023]; we treat that claim as the hypothesis under test and its lineage as the object we reproduce at parity. Two of our arms — a wired CfC (A_NCP) and a dense CfC (A_CFC) — are chosen precisely to isolate the contribution of the wiring from that of the continuous dynamics.

Continuous-time and state-space cousins. Deep state-space models share the continuous-time framing: S4 parameterizes long-range dependencies through a structured state-space kernel [Gu et al., 2022], Mamba adds input-dependent selectivity for linear-time sequence modeling [Gu and Dao, 2023], and Liquid-S4 connects the liquid formulation to this family [Hasani et al., 2023]. These situate our small recurrences within a broader continuous-time lineage; we do not evaluate them here.

Pre-registration, reproducibility, and imitation learning. Our methodology adopts the discipline of freezing success criteria before measurement and reporting honest baselines, in the spirit of the machine-learning reproducibility and pre-registration literature [Pineau et al., 2021, Søgaard et al., 2023]. Training uses behavioral cloning followed by DAgger, the canonical no-regret reduction of imitation learning to online learning [Ross et al., 2011]; a central engineering finding of this work concerns the interaction between the DAgger update mode and continuous-time cores (Section 6). The vision harness is built on gym-pybullet-drones [Panerati et al., 2021], and the companion open-loop regime uses the PASTIS satellite time-series benchmark [Sainte Fare Garnot and Landrieu, 2021].

3 The Measurement Program

The program is a house methodology applied uniformly across regimes.

Honest twin. Each arm shares the same encoder, the same `Linear(64→6)` head with identical output scaling, the same training data and budget; cores are matched to a reference of 27,648 core parameters to within $\pm 2\%$. The only sanctioned inter-arm difference is the learning rate, whose provenance is documented.

Frozen criteria; MEASURE = REPORT. Success criteria are frozen in a gate document before any measurement; whatever is measured is reported, including nulls and boundaries. The vision-loop gate (F3-GATE) fixes the OOD level (T2b), the primary metric, the decision threshold, and n before a single OOD scene is rendered.

Gates and preconditions. A capability **precondition** (nominal success $\geq 90\%$ per arm) must pass before the thesis metric is evaluated; the precondition is *not* the thesis. Failing it triggers a sanctioned, nominal-only repair (a learning rate from the pre-registered fallback grid — nothing else) and a full retrain, logged; exhaustion is a STOP for a human decision, not a verdict on the thesis.

Annexes with a stop rule. When the recipe diverges from a frozen, known-good reference, we unfreeze via an **annex** that cites the reference line-by-line and applies the change symmetrically to all arms. The stop rule (Annex-4) declares the last instrumental annex of phase 3a: if a CfC arm still fails the precondition afterward, FAIL becomes a reported **boundary**, not another annex.

Arithmetic early resolution. Because the precondition is a mean over 10 seeds, a leg can be resolved before all seeds run: after k seeds summing to S , if $S + (10 - k) \cdot 100 < 900$ the leg cannot reach 90% and is declared FAIL. This is arithmetic on the frozen criterion, not a change to it. In phase 3 it reduced the run to **13 full cycles instead of 40** and left one four-seed batch never launched (Appendix D).

Determinism contract. Environment, rendering, and seeding are bit-exact reproducible **within a single machine** (two runs, rgb/kin/setpoint identical). Cross-machine hashes may differ and that is not a failure — determinism is a within-machine property, not a cross-machine guarantee.

4 Harnesses

The portfolio uses three instruments, one per regime.

4.1 PASTIS twin (P0). Sentinel-2 crop classification of a phenologically confusable pair (soft winter wheat vs. corn) under test-time observation dropout, with a shared CNN encoder (28,752 parameters) and only the recurrent core swapped between CfC and GRU. Open-loop, day-scale. Companion; numbers verified against the companion report [Żydziać, 2026].

4.2 LiquidFlight (state-loop, v1.0). A closed-loop drone flight twin (gym-pybullet-drones [Panerati et al., 2021]) with a setpoint→DSL-PID inner loop at 48Hz, an observation-dropout OOD axis and a gap-length (latency) axis, and a time-constant (τ) inspection panel; CfC(32) vs. GRU at core parity, everything but the core bit-identical between arms. The vision-loop harness reuses this execution layer verbatim (sha256-pinned copies).

4.3 LiquidSight (vision-loop, phase 3 — this work). A fly-to-target task from a $64 \times 64 \times 3$ forward camera plus kinematics: reach a ground marker and hold a hover within `r_goal` for the final dwell window. Physics at 240 Hz, control at 48 Hz, camera and policy tick at 12 Hz; the target lives only in pixels. A privileged DSL-PID expert supplies behavioral cloning (BC) labels; three DAgger rounds follow. The perceptual-shift axis is a distractor ladder T0–T3 (T2/T2a/T2b/T2c = $K \in \{0, 1, 2, 3\}$ distractors on structured backgrounds), with P-SANITY establishing that the axis resolves difficulty and that the expert ceiling is reachable at every level (Section 7). The gate level is T2b.

5 Results Across Regimes

The boundary map (Table 1) summarizes the direction and measurement resolution of the effect in each regime. No measurement crosses its pre-registered threshold.

Table 1: Boundary map across task regimes.

row	regime / harness	effect direction	verdict / resolution
R1	open-loop classification (P0)	weak positive under dropout	null (margin < pooled std); quantitative
R2	onset detection, low FAR (E6)	negative	directional; qualitative (secondary source)
R3	state-loop control (v1.0)	zero	quantitative
R4	vision-loop (this work)	untestable under precondition	precondition FAIL at parity; quantitative

5.1 Vision-loop (phase 3) — the primary result

5.1 The robustness claim is untestable under our precondition (C1). After restoring the known-good recipe in full (core construction, Annex-3; training procedure, Annex-4) and exhausting the only sanctioned lever — the learning-rate grid $\{3e-4, 1e-3\}$ — **no CfC arm reaches the pre-registered precondition** of mean $\geq 90\%$ nominal success at core parity. All four lr legs resolve to FAIL by arithmetic early resolution: A_NCP at $3e-4 = 16.5\%$ (resolved at $k=2$), A_NCP at $1e-3 = 72.2\%$ ($k=4$), A_CFC at $3e-4 = 27.0\%$ ($k=2$), A_CFC at $1e-3 = 55.3\%$ ($k=3$). Because the **deciding arm** (A_NCP , the AutoNCP-wired CfC faithful to the wiring in the cited claim) does not qualify, the thesis metric — a T2b margin of A_NCP over A_GRU with threshold $M > \text{pooled_std}$, $n=10$ — is *untestable under our precondition*. We therefore ran **zero OOD evaluations** in the entire phase: gate stages B (the A_GRU precondition sweep) and C (the OOD ladder) were skipped by the branch, so the pre-registered verdict remains untouched and unprejudiced.

This is deliberately *not* a falsification of the thesis. A test requires a qualifying deciding arm; without the precondition there is no test. The result is a cell in the boundary map, not a verdict on the mechanism.




Figure 1: *Per-seed nominal success for each lr leg.* Points show nominal success (100 scenes, level T0) for every trained seed across the four CfC lr legs (A_NCP and $A_CFC \times \{3e-4, 1e-3\}$); horizontal bars mark per-leg means; a dashed rule marks the 90% precondition. The A_GRU control (100%) is shown for reference. The wired arm at 1e-3 spans 49–92% (43 pp); no leg mean crosses 90%.

5.2 A trainability tax at parity (C2). At identical data, budget, and procedure, and with cores matched to $\pm 2\%$, nominal capability forms a gradient: **A_GRU 100%** (stable; dwell failures 0) > **wired CfC** (A_NCP ; peak 92% at seed 45010/1e-3; per-leg mean 72.2% at $k=4$) > **dense CfC** (A_CFC ; cap 65% at seed 45012/1e-3). Core-parameter parity alone does not buy capability; the continuous-time core pays a trainability tax at this budget. The dominant residual failure is **dwell** — the drone reaches the marker but does not hold the hover precisely — while catastrophic **tilt is eliminated** at 1e-3 (0 tilt across all CfC cycles, versus 78/100 for the dense CfC under the pre-Annex-4 procedure).

5.3 The boundary is variance, not mean (C3). The wired arm at 1e-3 spans **49–92% across seeds** — a **43 pp spread** — under one fixed procedure, while the GRU control is a stable 100%. Reporting only the $k=4$ mean (72.2%) understates the phenomenon; the **full six-seed characterization (75.3%, spread 43 pp)** shows the boundary is a property of the between-seed distribution, not a sample-size artifact. The CfC at parity does not “fail to learn”; it learns **unstably and dwell-limited, on average below 90%**. This mirrors the companion P0 result [Żydziać, 2026], where the retention margin (+0.1127) and its pooled standard deviation (0.1759) are nearly invariant between $n=3$ and $n=15$ — population spread that does not shrink with n .

5.4 Wiring attribution (C4). The contrast between arms attributes trainability to the wiring: the AutoNCP-wired CfC reaches a 92% peak and per-leg means up to 75.3%, whereas the dense CfC caps at 65%, across seeds and at the 1e-3 lever. The attribution holds on the 1e-3 lever; at 3e-4 the ordering inverts (dense 27.0% > wired 16.5%), so the wiring benefit is itself learning-rate dependent and we report it as such. A historical, single-seed diagnostic corroborates the readout dimension: reading the full state rather than six motor neurons raises reach from 8/50 to 26/50

(Section 6).




Figure 2: *DAgger rollout dynamics per round.* On-policy rollout success across DAgger rounds $r1 \rightarrow r2 \rightarrow r3$. The GRU control climbs $18 \rightarrow 100 \rightarrow 100$; CfC arms stay low or flat, illustrating that the aggregate does not close for the continuous-time cores (thin: per-seed; thick: per-leg mean).

5.2 Open-loop classification, day-scale (P0) — companion

The PASTIS twin gives a directional-but-null signal: at full cadence the GRU leads on macro-F1 (0.9073 ± 0.0137 vs. 0.8802 ± 0.0121), but under observation dropout the CfC retains more — retention $R(0.6)$ of 0.6415 ± 0.1138 vs. 0.5288 ± 0.0622 , a **+0.1127 margin against a pooled std of 0.1759, i.e. null** by the pre-registered rule. An absolute-F1 crossover appears at $d=0.6$ (CfC 0.5642 ± 0.0985 above GRU 0.4798 ± 0.0567 , despite a -0.027 start), and the degradation slope is gentler for the CfC (-0.5095 vs. -0.6936). Crucially, margin and pooled std are near-invariant from $n=3$ to $n=15$ — the boundary is population spread, not sampling noise. All P0 numbers are verified against the companion report [Żydzia, 2026].

5.3 Open-loop onset detection, low FAR (E6) — companion

The onset-detection regime contributes the **negative direction** of the map: at a low false-alarm-rate operating point the effect sign is **unfavorable to the liquid core**, as recorded in the program compendium; the primary E6 report is not preserved in the project archive. We therefore state this regime **qualitatively only** and report no numerical values for it — in particular the detection-delay delta, the sweep parameter, and the false-alarm-rate definition are not carried into the prose, since their primary source could not be verified. The regime enters the boundary map as a directional cell (Table 1, row R2), not as a quantitative result.

5.4 State-loop control, millisecond-scale (LiquidFlight v1.0) — companion

In the state-loop regime the continuous-time lever (Δt) is not behaviorally load-bearing. With a setpoint action abstraction the CfC-32 policy flies under 500–1300 ms observation gaps and the

episode-failure cliff moves from ≈ 102 ms (raw-RPM stabilization) to ≈ 779 ms; the Δt mechanism is *visible* in the state dynamics (recovery after a gap) but yields **no behavioral edge** — ablating the Δt channel flies identically. The fitted time constants are short (median $\tau \approx 35$ ms, IQR 24–69 ms), so a 500 ms gap is $\approx 14 \times \tau$ and the model does not bridge gaps with long memory; τ serves as an explainability handle, not a source of advantage. Under the A1 perturbation grid there were **0 tip-overs**. The regime contributes a **zero** to the map: neither help nor harm attributable to the liquid core at this horizon.

6 Engineering Findings

Bringing a CfC to the point where it will run at all in this harness required restoring a known-good recipe in four provenanced steps (fixes **F1–F4**); each is an **engineering finding, established by single-seed controlled probes (diagnostic, not statistical)** — not a statistically powered ablation.

F1 — Backbone. The dense CfC built without a backbone under the parity constraint does not reach the target (reach 8/50); adding a frozen-style backbone restores it (35/50, $\approx$ GRU) (Annex-3).

F2 — The Δt path and its units (C5). The continuous-time core — the sole locus of the thesis mechanism — is load-bearing only after two corrections. First, the time-span scale sets the gate regime: with an explicit **ts** in **seconds** (0.0833 s per 12 Hz tick) reach is 39/50, versus 15/50 at $ts=1.0$ (ticks) and 9/50 at $ts=4.0$; the post-init gate drive $|t_a \cdot ts|$ scales accordingly (0.010 / 0.125 / 0.501), so a mis-scaled **ts** freezes or destabilizes the cell. Second, we found a **library bug**: `ncps.torch.CfC` rejects explicit **timespans** when the batch size exceeds one (scalar, $(B, 1)$, and $(B,)$ all raise; only `None=1.0` works), which had trapped the arms at $ts=1.0$. We work around it by stepping the cell manually. The auxiliary thesis is therefore: **Δt is load-bearing implementationally, not advantageously** — in this harness we could not measure an advantage because the precondition was never met.

F3 — Full-state readout. Reading the wired core’s full state rather than six motor neurons lowers BC error (0.0047 vs. 0.0220) and raises reach (26/50 vs. 8/50) (Annex-3).

F4 — DAgger update mode (C6). Continuous-time cells are sensitive to the DAgger update procedure. Under the gate’s original v1 mode (continue from the previous weights + final-epoch checkpoint), the dense CfC becomes unstable (tilt 78/100) and the wired CfC collapses its rollouts (1 $\rightarrow$ 0 $\rightarrow$ 0); the dense CfC is stable after BC alone (0 tilt), so the instability is **induced by the procedure**, not the construction. Adopting the frozen-C1 mode symmetrically (retrain from scratch each round + best-validation checkpoint, 120 epochs per stage; Annex-4) removes the tilt (78 $\rightarrow$ 0 at $1e-3$) and follows the canonical DAgger formulation [Ross et al., 2011]. The **GRU is robust to both procedures** (100% either way).

A mechanistic signature accompanies this (C6a): for every CfC arm the best validation error after BC alone is low (≈ 0.00012 – 0.0009) but **rises** after the DAgger aggregate (rounds r1–r3), and the best epoch moves early (e.g., wired at $1e-3$, seed 45013: 119 $\rightarrow$ 41 $\rightarrow$ 26 $\rightarrow$ 37) — the core does not close the fit to the policy-visited distribution as well as it fits the clean expert. The GRU holds its best validation low across all rounds (.000168 $\rightarrow$.000112) — a different regime.




Figure 3: *Best-validation trend, BC vs. aggregate.* Best validation MSE per stage ($r_0 = \text{BC} \rightarrow r_1\text{--}r_3 = \text{Dagger aggregate}$) for the CfC arms (rising after aggregation, best epoch moving earlier) against the GRU control (monotonically low). Round-level (four points per cycle); full per-epoch curves are not persisted and would require re-measurement.

7 Instrument Characterization and Protocol Economics

Axis resolution and expert ceiling (C8). The perceptual-shift axis resolves difficulty monotonically. The sanity policy degrades **100/100/64/46/36/24/16** across T0/T1/T2/T2a/T2b/T2c/T3 (50 scenes per level), placing three levels inside the pre-registered [30%, 85%] band and selecting **T2b (36%)** as the hardest in-band gate level. The privileged expert reaches a **100% ceiling at every level**, confirming that the distractors do not block physical reachability — the axis touches pixels only. Determinism is bit-exact **within a machine** (rgb/kin/setpoint identical across two runs; cross-machine equality is neither required nor claimed).

Protocol economics (C7a). The binding run comprised **13 full training cycles** (plus one GRU control), **≈57.1 h** of cumulative compute (sum of per-cycle wall time; parallel wall-clock was shorter but is not an authoritative per-arm cost). Arithmetic early resolution spared a complete four-seed batch relative to the 40 cycles a full $n=10$ over four legs would have required.

8 Discussion

Synthesis. Across the regimes we map, the sign of the effect depends on the task: a weak positive in sparse open-loop observation (P0), a negative under low-FAR onset detection (E6), a zero in state-loop control (LiquidFlight), and — in the richest regime, vision-in-the-loop — a question that never reaches the conditions of measurement, because the advantage is preceded by a trainability tax at parity. No measurement crosses its pre-registered threshold. The unifying reading is that **task construction decides whether the mechanism can reveal itself at all**: the continuous-time core needs a regime whose dependencies it is positioned to exploit, and at parity it must first be trainable there.




Figure 4: *Sanity ladder and expert ceiling.* Policy success across the distractor ladder (T0→T3: 100/100/64/46/36/24/16) with the [30,85] band shaded and the T2b gate level marked; the expert ceiling (100% at every level) overlaid.

Demonstration culture vs. falsification. A pre-registered null carries evidential weight that a single favorable curve does not. Much of the liquid-network robustness literature is demonstrated rather than falsified: a favorable model is shown on a task chosen after the fact, without a same-budget parity baseline and without a criterion fixed in advance. Our program inverts that stance — it commits to a threshold, a level, and a seed count before any out-of-distribution scene is seen, and it reports whatever the instrument returns, including a boundary. The cost of this discipline is that we sometimes conclude “untestable” rather than “wins” or “loses”; the benefit is that the three cells we *can* state (weak-positive, negative, zero) are not artifacts of selective reporting. The map says where to look and where not to — it does not claim the mechanism is absent, only that under parity, fixed budget, and frozen criteria we did not measure it in these harnesses.

When the twin re-arms. The vision-loop cell would re-enter measurement under conditions that unlock trainability without breaking parity. The cleanest such condition is a **symmetric budget increase** for all arms — more epochs, more DAgger rounds, or more data applied identically to GRU and CfC — carried until the deciding arm (A_NCP) meets the precondition; the wired arm at 1e-3 is the natural starting point, already at 92% best seed and 75.3% mean. Because the increase is symmetric, the duel remains fair: if the GRU also improves, the parity contract is preserved. Two subsidiary questions gate the value of that rerun. First, the sensitivity of continuous-time cells to the DAgger update mode (F4) is an open research problem worth isolating on its own — why does retrain-from-scratch with best-validation selection stabilize a CfC that continuation destabilizes, when a GRU is indifferent to both? Second, whether the residual dwell failure reflects a precision limit of the continuous-time state under this readout, or merely under-training, is answerable only once the precondition is met. Until then, we resist reading the boundary as evidence about the mechanism itself.

Next habitat. The mechanism’s most plausible next habitat is a regime with genuinely variable time steps or asynchronous/event-stream sensing, where the continuous-time gate has something to integrate — offered as a direction, not a promise.

9 Limitations

- **Simulation only.** PyBullet dynamics, no photorealism, no sim-to-real transfer.
- **One harness per regime.** A single fly-to-target task family at 64×64 ; results may not generalize across task families.
- **Parity at $\approx 27k$ core parameters.** Findings are stated at this capacity; other capacities are untested.
- **$n=10$ with a finality rule.** Seeds 45010–45015 were used (early resolution); increasing n , or changing the level or threshold after seeing any OOD result, is forbidden by the gate.
- **No OOD measurement in the vision-loop regime.** This is a *consequence* of the unmet precondition, not a design choice.
- **Engineering findings are single-seed probes.** The four recipe restorations (fixes F1–F4) are diagnostic controlled probes ($n=1$ for most), illustrative rather than statistically powered.
- **Determinism is within-machine,** not cross-machine.
- **Fixed training budget.** The trainability tax is measured at BC-120 + DAgger $\times 3$; a higher symmetric budget is unmeasured.
- **Single-operator program;** one task family per regime.
- **Companion E6 primary report not preserved.** The onset-detection regime (R2) is stated qualitatively from a secondary source; its primary report is absent from the archive, so R2 carries no numerical values and its direction should be read as directional evidence, not a measured effect size.

10 Reproducibility and Artifacts

We release the liquidflight v1.0 and liquidsight repositories (tags and gate commits), the frozen documents (F3_PRE0, decisions and Annexes 1–4, P-SANITY, F3-GATE), sha256 manifests of the execution layer, the seed pools, the per-cycle logs, and the reports. The companion P0 ships as a separate PDF/arXiv artifact [Żydziać, 2026]. Instrument determinism is within-machine bit-exact; bit-for-bit reproduction across different GPUs is not guaranteed (inherent to CUDA).

11 Conclusion

We set out to test a robustness claim about liquid networks and, in the richest regime, found the claim **untestable under our pre-registered precondition**: at core parity and fixed budget, no continuous-time arm becomes reliably competent, so no out-of-distribution verdict can be earned. The first-order result is therefore a **boundary map** — where the mechanism helps (weakly), hurts, is neutral, and cannot yet be measured — together with a **trainability tax** that the field

does not usually report: a same-budget capability gradient from a stable GRU (100%) through a wired CfC (peak 92%) to a dense CfC (cap 65%). The measurement methodology — gates, preconditions, provenanced annexes with a stop rule, and arithmetic early resolution — is the portable contribution. A negative result, pre-registered, is a map of where to measure next.

A Program timeline




Figure 5: *Gate and annex timeline.* The phase-3 program as a left-to-right order (2026-07-22 to 07-24): frozen gates (P-SANITY, F3-GATE, parity v1/v2), annexes (A1 observability, A2 distractor ladder, A3 core construction, A4 training procedure), and measurements (I3a, DIAG, I3a-R, I3a-R2, I3b).

One-line reasons: **Annex-1** spawned the target in the frontal cone and pitched the camera -22.3° so the marker is in-frame; **Annex-2** extended the axis with T2a/T2b/T2c ($K=1/2/3$) to fill resolution; **Annex-3** fixed three CfC construction defects (backbone, ts-in-seconds, full-state readout); **Annex-4** unified the training procedure with frozen C1 (from-scratch per round, best-val, 120 epochs) and declared the stop rule.

B Frozen criteria (pointers)

- **P-SANITY:** P1 capability $\geq 90\%$; P2 axis resolution ≥ 2 levels in $[30, 85]$; P3 expert ceiling $\geq 95\%$ /level.
- **F3-GATE §4:** precondition mean nominal success $\geq 90\%$ per arm; fallback grid $\{3e-4, 1e-3\}$; no OOD before all arms pass.
- **F3-GATE §5:** primary metric $M = \text{mean}_s \text{succ}(A_NCP, s) - \text{mean}_s \text{succ}(A_GRU, s)$, $n=10$; $\text{pooled_std} = \sqrt{(\text{sd}(A_NCP)^2 + \text{sd}(A_GRU)^2)/2}$; verdict $M > \text{pooled_std}$; binding at $n=10$.
- **Stop rule** (Annex-4): the last instrumental annex of phase 3a; a subsequent precondition FAIL becomes a reported boundary.

C Per-seed tables (binding run I3b)

Columns: nominal % | failures | DAgger rollout r1→r2→r3 | best-val r0→r3 | best-epoch | cycle seconds. Generated from the per-cycle log (read, not measured).

Table 2: A_NCP @ 1e-3 (best leg; mean 72.2% at $k=4$, 75.3% over six seeds, 43pp spread).

seed	nom%	failures	rollout	best-val r0→r3	sec
45010	92	dwell 8	44→32→53	.000174→.000587	18,889.3
45011	79	dwell 21	43→37→52	.000153→.000696	19,138.1
45012	69	dwell 31	25→11→36	.000169→.001187	19,441.2
45013	49	dwell 51	9→1→13	.000184→.001711	18,986.6
45014	86	dwell 14	32→12→37	.000121→.000878	19,471.4
45015	77	dwell 23	25→38→23	.000152→.000934	18,992.2

Table 3: Other legs and the GRU control (procedure v2).

leg	seed	nom%	failures	rollout	best-val r0→r3
A_NCP 3e-4	45010	22	dwell 76, tilt 2	22→13→25	.000566→.001843
A_NCP 3e-4	45011	11	dwell 89	12→0→12	.000721→.002428
A_CFC 1e-3	45010	44	dwell 56	0→8→44	.000648→.000965
A_CFC 1e-3	45011	57	dwell 43	2→14→40	.000631→.001127
A_CFC 1e-3	45012	65	dwell 35	0→34→44	.00047→.001041
A_CFC 3e-4	45010	18	dwell 70, tilt 12	8→0→6	.000887→.002938
A_CFC 3e-4	45011	36	dwell 64	2→0→8	.000889→.002709
A_GRU 1e-3	45010	100	— (dwell 0)	18→100→100	.000168→.000112

D Early-resolution formula and savings

For a per-arm precondition of mean $\geq 90\%$ over 10 seeds, after k seeds with partial sum S the leg is declared FAIL when

$$S + (10 - k) \cdot 100 < 900.$$

Applied to the four legs: A_NCP@3e-4 FAIL at $k=2$ ($S=33$; $833 < 900$); A_NCP@1e-3 FAIL at $k=4$ ($S=289$; $889 < 900$); A_CFC@3e-4 FAIL at $k=2$ ($S=54$; $854 < 900$); A_CFC@1e-3 FAIL at $k=3$ ($S=166$; $866 < 900$). Total executed: 13 cycles versus 40 for a full $n=10$ over four legs; one four-seed batch (45016–45019) was never launched. This is arithmetic on the frozen criterion, not a modification of it.

References

Makram Chahine, Ramin Hasani, Patrick Kao, Aaron Ray, Ryan Shubert, Mathias Lechner, Alexander Amini, and Daniela Rus. Robust flight navigation out of distribution with liquid neural networks. *Science Robotics*, 2023.

Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces, 2023.

- Albert Gu, Karan Goel, and Christopher Ré. Efficiently modeling long sequences with structured state spaces. In *International Conference on Learning Representations (ICLR)*, 2022. arXiv:2111.00396.
- Ramin Hasani, Mathias Lechner, Alexander Amini, Daniela Rus, and Radu Grosu. Liquid time-constant networks. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2021.
- Ramin Hasani, Mathias Lechner, Alexander Amini, Lucas Liebenwein, Aaron Ray, Max Tschaikowski, Gerald Teschl, and Daniela Rus. Closed-form continuous-time neural networks. *Nature Machine Intelligence*, 2022.
- Ramin Hasani, Mathias Lechner, Tsun-Hsuan Wang, Makram Chahine, Alexander Amini, and Daniela Rus. Liquid structural state-space models. In *International Conference on Learning Representations (ICLR)*, 2023. arXiv:2209.12951.
- Mathias Lechner, Ramin Hasani, Alexander Amini, Thomas A. Henzinger, Daniela Rus, and Radu Grosu. Neural circuit policies enabling auditable autonomy. *Nature Machine Intelligence*, 2: 642–652, 2020. doi: 10.1038/s42256-020-00237-3.
- Jacopo Panerati, Hehui Zheng, SiQi Zhou, James Xu, Amanda Prorok, and Angela P. Schoellig. Learning to fly—a gym environment with pybullet physics for reinforcement learning of multi-agent quadcopter control. In *2021 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, page [BIB:verify], 2021. arXiv:2103.02142.
- Joelle Pineau, Philippe Vincent-Lamarre, Koustuv Sinha, Vincent Larivière, Alina Beygelzimer, Florence d’Alché Buc, Emily Fox, and Hugo Larochelle. Improving reproducibility in machine learning research (a report from the neurips 2019 reproducibility program). *Journal of Machine Learning Research*, 22(164):1–20, 2021.
- Stéphane Ross, Geoffrey J. Gordon, and J. Andrew Bagnell. A reduction of imitation learning and structured prediction to no-regret online learning. In *Proceedings of the 14th International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 15 of *Proceedings of Machine Learning Research*, pages 627–635, 2011.
- Vivien Sainte Fare Garnot and Loic Landrieu. Panoptic segmentation of satellite image time series with convolutional temporal attention networks. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021.
- Anders Søgaard, Daniel Hershcovich, and Miryam de Lhoneux. A two-sided discussion of preregistration of NLP research. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, 2023. ACL Anthology 2023.eacl-main.6; arXiv:2302.10086.
- Olga Żydzia. Temporal robustness of liquid neural networks under irregular observation dropout: A controlled null result on satellite image time series. Technical report, Independent Research, July 2026.